Harsh Banerjee

harsh.banerjee@example.com | +91 99569 19793 | GitHub | LinkedIn | Portfolio

PROFESSIONAL SUMMARY

Results-driven software engineer with strong foundation in distributed systems, high-concurrency backend services, and clean architecture. Experienced in designing fault-tolerant microservices, optimizing database throughput, and shipping production code with rigorous test coverage.

EDUCATION

Master of Technology in Software Systems — Birla Institute of Technology and Science *2022 - 2026*

Cumulative Grade Point Average: 8.40/10.0

WORK EXPERIENCE

Associate Software Engineer — Swiggy (Bundl Technologies), Bengaluru, India *Jun 2025 - Present*

- Migrated monolithic checkout service to gRPC microservices, shrinking cloud compute costs by 32% and reducing deployment cycle times from 4 hours to 12 minutes.
- Developed responsive internal operations dashboard in React and TypeScript with TanStack Query, accelerating merchant dispute resolution by 40%.
- Designed real-time event pipeline consuming 5M daily telemetry records via Apache Kafka, providing sub-second fraud detection alerts.

TECHNICAL PROJECTS

Autonomous Code Review Linter & Security Auditor | Technologies: Python, AST, Tree-sitter, Docker, FastAPI

- Engineered static analysis pipeline parsing abstract syntax trees (ASTs) to detect SQL injection, unvalidated deserialization, and memory leaks.
- Built automated GitHub Action integration running checks across 200+ pull requests with an average scan duration under 3.5 seconds.
- Containerized execution sandbox using gVisor and Docker to securely isolate untrusted student code submissions.

Distributed Log-Structured Key-Value Store | Technologies: Go, Raft Consensus, gRPC, LSM-Tree, Docker

- Implemented distributed LSM-tree storage engine in Go with write-ahead logging (WAL), SSTable compaction, and bloom filters.
- Embedded Raft consensus algorithm for distributed multi-node replication, surviving arbitrary node crashes without state divergence.
- Benchmarked performance using custom concurrent load generator, achieving 85,000 read operations/sec across a 5-node cluster.

TECHNICAL SKILLS

Languages: Go, Java, Python, TypeScript, C++, SQL

Frameworks & Systems: Spring Boot, Node.js, FastAPI, React, Next.js, gRPC

Databases & Storage: PostgreSQL, Redis, MongoDB, Apache Kafka, Elasticsearch

Cloud & DevOps: Docker, Kubernetes, AWS (EC2, S3, RDS), Git, CI/CD Pipelines